

Name :- Abhijeet Dinakar Pathade

Roll_No. :- 16

Division :- A

✓ KNN Classification on the Pima Indians Diabetes Dataset

This notebook builds a **K-Nearest Neighbors (KNN)** classifier that predicts whether a patient has diabetes based on diagnostic measurements.

Workflow

1. Load and inspect the data
2. Clean invalid (zero-coded) medical readings
3. Explore the data visually
4. Split and scale the features
5. Tune the hyperparameter **K** using cross-validation
6. Train the final model and evaluate it (accuracy, confusion matrix, ROC-AUC)
7. Summarize findings

Dataset: Pima Indians Diabetes Database — 768 female patients, 8 medical predictor variables, and a binary target **Outcome** (1 = diabetic, 0 = not diabetic).

✓ 1. Setup & Imports

```
# Core libraries
import numpy as np
import pandas as pd

# Visualization
import matplotlib.pyplot as plt
import seaborn as sns
sns.set_style("whitegrid")
plt.rcParams["figure.figsize"] = (8, 5)

# Modeling
from sklearn.model_selection import train_test_split, cross_val_score, GridSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import (
    accuracy_score, confusion_matrix, classification_report,
    ConfusionMatrixDisplay, roc_curve, roc_auc_score
)

RANDOM_STATE = 42
print("Libraries loaded.")
```

Libraries loaded.

✓ 2. Load the Dataset

The dataset is fetched directly from a public GitHub mirror of the Pima Indians Diabetes Database, so this cell will work as-is in Google Colab (internet access required).

If the URL ever becomes unavailable, you can instead:

- Download **diabetes.csv** from [Kaggle: uciml/pima-indians-diabetes-database](https://www.kaggle.com/uciml/pima-indians-diabetes-database), then
- Upload it in Colab with the file-upload cell provided (commented out below).

```
url = "/content/drive/MyDrive/diabetes_dataset.csv"
df = pd.read_csv(url)

# --- Alternative: upload the CSV manually in Colab ---
# from google.colab import files
```

```
# uploaded = files.upload()
# df = pd.read_csv(list(uploaded.keys())[0])

print("Shape:", df.shape)
df.head()
```

Shape: (768, 9)

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1

3. Exploratory Data Analysis

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Pregnancies            768 non-null   int64
1   Glucose                768 non-null   int64
2   BloodPressure          768 non-null   int64
3   SkinThickness          768 non-null   int64
4   Insulin                768 non-null   int64
5   BMI                    768 non-null   float64
6   DiabetesPedigreeFunction 768 non-null   float64
7   Age                    768 non-null   int64
8   Outcome                768 non-null   int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

```
df.describe().T
```

	count	mean	std	min	25%	50%	75%	max
Pregnancies	768.0	3.845052	3.369578	0.000	1.00000	3.0000	6.00000	17.00
Glucose	768.0	120.894531	31.972618	0.000	99.00000	117.0000	140.25000	199.00
BloodPressure	768.0	69.105469	19.355807	0.000	62.00000	72.0000	80.00000	122.00
SkinThickness	768.0	20.536458	15.952218	0.000	0.00000	23.0000	32.00000	99.00
Insulin	768.0	79.799479	115.244002	0.000	0.00000	30.5000	127.25000	846.00
BMI	768.0	31.992578	7.884160	0.000	27.30000	32.0000	36.60000	67.10
DiabetesPedigreeFunction	768.0	0.471876	0.331329	0.078	0.24375	0.3725	0.62625	2.42
Age	768.0	33.240885	11.760232	21.000	24.00000	29.0000	41.00000	81.00
Outcome	768.0	0.348958	0.476951	0.000	0.00000	0.0000	1.00000	1.00


Class balance

Diabetes datasets are commonly imbalanced. Let's check how many patients are diabetic (`Outcome = 1`) vs not (`Outcome = 0`).

```
print(df["Outcome"].value_counts())
print(df["Outcome"].value_counts(normalize=True).round(3) * 100, "%")

sns.countplot(x="Outcome", hue="Outcome", data=df, palette=["#3E5C76", "#8C4A4A"], legend=False)
plt.title("Class Distribution (0 = No Diabetes, 1 = Diabetes)")
plt.xlabel("Outcome")
plt.ylabel("Number of Patients")
plt.show()
```

```
Outcome
0    500
1    268
Name: count, dtype: int64
Outcome
0    65.1
1    34.9
Name: proportion, dtype: float64 %
```



✓ Hidden missing values

Several columns cannot biologically be 0 (e.g. Glucose, BloodPressure, BMI). In this dataset those missing readings were encoded as 0. Let's find them before cleaning.

```
zero_invalid_cols = ["Glucose", "BloodPressure", "SkinThickness", "Insulin", "BMI"]

zero_counts = (df[zero_invalid_cols] == 0).sum()
print("Number of zero (likely missing) values per column:")
print(zero_counts)
```

```
Number of zero (likely missing) values per column:
Glucose      5
BloodPressure 35
SkinThickness 227
Insulin      374
BMI          11
dtype: int64
```

✓ 4. Data Cleaning

We replace the biologically-impossible zeros with `NaN`, then impute them using the **median of each column, computed separately per class (Outcome)**. This keeps the imputed values consistent with whether the patient is diabetic or not, which is more informative than a single global median.

```
df_clean = df.copy()
df_clean[zero_invalid_cols] = df_clean[zero_invalid_cols].replace(0, np.nan)

# Impute with the median of each column, computed per Outcome class
for col in zero_invalid_cols:
    df_clean[col] = df_clean.groupby("Outcome")[col].transform(lambda s: s.fillna(s.median()))

print("Remaining missing values:\n", df_clean.isna().sum().sum())
df_clean.describe().T
```


Remaining missing values:
0

	count	mean	std	min	25%	50%	75%	max
Pregnancies	768.0	3.845052	3.369578	0.000	1.00000	3.0000	6.00000	17.00
Glucose	768.0	121.677083	30.464161	44.000	99.75000	117.0000	140.25000	199.00
BloodPressure	768.0	72.389323	12.106039	24.000	64.00000	72.0000	80.00000	122.00
SkinThickness	768.0	29.089844	8.890820	7.000	25.00000	28.0000	32.00000	99.00
Insulin	768.0	141.753906	89.100847	14.000	102.50000	102.5000	169.50000	846.00
BMI	768.0	32.434635	6.880498	18.200	27.50000	32.0500	36.60000	67.10
DiabetesPedigreeFunction	768.0	0.471876	0.331329	0.078	0.24375	0.3725	0.62625	2.42
Age	768.0	33.240885	11.760232	21.000	24.00000	29.0000	41.00000	81.00
Outcome	768.0	0.348958	0.476951	0.000	0.00000	0.0000	1.00000	1.00


5. Visualizing the Data

```
df_clean.hist(figsize=(12, 10), bins=20, color="#3E5C76", edgecolor="white")
plt.suptitle("Feature Distributions (after cleaning)", y=1.02, fontsize=14)
plt.tight_layout()
plt.show()
```

Feature Distributions (after cleaning)




```
plt.figure(figsize=(9, 7))
corr = df_clean.corr()
sns.heatmap(corr, annot=True, fmt=".2f", cmap="vlag", center=0, square=True, cbar_kws={"shrink": 0.8})
plt.title("Feature Correlation Matrix")
plt.show()
```



```
features_to_plot = ["Glucose", "BMI", "Age", "Insulin"]

fig, axes = plt.subplots(1, len(features_to_plot), figsize=(16, 4))
for ax, feat in zip(axes, features_to_plot):
    sns.boxplot(x="Outcome", y=feat, hue="Outcome", data=df_clean, ax=ax, palette=["#3E5C76", "#8C4A4A"], legend=False)
    ax.set_title(feat)
plt.tight_layout()
plt.show()
```



6. Train/Test Split & Feature Scaling

KNN is a **distance-based** algorithm, so features must be on comparable scales — otherwise a large-range feature like `Insulin` would dominate the distance calculation over a small-range feature like `DiabetesPedigreeFunction`. We standardize all features to zero mean and unit variance using `StandardScaler`, fit **only on the training set** to avoid data leakage.

```
X = df_clean.drop(columns=["Outcome"])
y = df_clean["Outcome"]
```

```
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=RANDOM_STATE, stratify=y
)

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

print("Train shape:", X_train_scaled.shape)
print("Test shape :", X_test_scaled.shape)
```

```
Train shape: (614, 8)
Test shape : (154, 8)
```

7. Choosing the Best K

We evaluate odd values of K from 1 to 30 using 10-fold cross-validation on the training set, then plot mean accuracy vs. K to pick the value that generalizes best (avoiding both overfitting at low K and underfitting at high K).


```
k_values = list(range(1, 31, 2)) # odd K avoids tie votes for binary classification
cv_scores = []

for k in k_values:
    knn = KNeighborsClassifier(n_neighbors=k)
    scores = cross_val_score(knn, X_train_scaled, y_train, cv=10, scoring="accuracy")
    cv_scores.append(scores.mean())

best_k = k_values[int(np.argmax(cv_scores))]
print(f"Best K found: {best_k} (CV accuracy = {max(cv_scores):.4f})")

plt.figure(figsize=(8, 5))
plt.plot(k_values, cv_scores, marker="o", color="#3E5C76")
plt.axvline(best_k, color="#8C4A4A", linestyle="--", label=f"Best K = {best_k}")
plt.title("Cross-Validated Accuracy vs. K")
plt.xlabel("Number of Neighbors (K)")
plt.ylabel("Mean CV Accuracy")
plt.legend()
plt.show()
```

Best K found: 17 (CV accuracy = 0.8242)



8. Train the Final KNN Classifier

```
knn_final = KNeighborsClassifier(n_neighbors=best_k, metric="minkowski", p=2) # p=2 -> Euclidean distance
knn_final.fit(X_train_scaled, y_train)

y_pred = knn_final.predict(X_test_scaled)
```

```
y_proba = knn_final.predict_proba(X_test_scaled)[: , 1]
```

```
print(f"KNN classifier trained with K = {best_k}")
```

```
KNN classifier trained with K = 17
```

9. Model Evaluation


```
acc = accuracy_score(y_test, y_pred)
print(f"Test Accuracy: {acc:.4f}\n")
print("Classification Report:\n")
print(classification_report(y_test, y_pred, target_names=["No Diabetes", "Diabetes"]))
```

```
Test Accuracy: 0.7857
```

```
Classification Report:
```

	precision	recall	f1-score	support
No Diabetes	0.83	0.85	0.84	100
Diabetes	0.71	0.67	0.69	54
accuracy			0.79	154
macro avg	0.77	0.76	0.76	154
weighted avg	0.78	0.79	0.78	154

```
cm = confusion_matrix(y_test, y_pred)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=["No Diabetes", "Diabetes"])
disp.plot(cmap="Blues", values_format="d")
plt.title(f"Confusion Matrix (K = {best_k})")
plt.grid(False)
plt.show()
```




ROC Curve

The ROC curve shows the trade-off between true-positive rate and false-positive rate across classification thresholds. AUC (Area Under the Curve) summarizes overall discriminative ability — 0.5 is random guessing, 1.0 is a perfect classifier.

```
fpr, tpr, _ = roc_curve(y_test, y_proba)
auc = roc_auc_score(y_test, y_proba)

plt.figure(figsize=(6.5, 6))
plt.plot(fpr, tpr, color="#3E5C76", label=f"KNN (AUC = {auc:.3f})")
plt.plot([0, 1], [0, 1], linestyle="--", color="gray", label="Random guess")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
```

```
plt.title("ROC Curve")
plt.legend()
plt.show()
```



10. Predicting a New Patient

Example of how to classify a brand-new, unseen patient record using the trained pipeline (scale, then predict).

```
new_patient = pd.DataFrame([{"Pregnancies": 2,
    "Glucose": 130,
    "BloodPressure": 70,
    "SkinThickness": 25,
    "Insulin": 100,
    "BMI": 31.5,
    "DiabetesPedigreeFunction": 0.45,
    "Age": 35}])

new_patient_scaled = scaler.transform(new_patient)
pred_class = knn_final.predict(new_patient_scaled)[0]
pred_proba = knn_final.predict_proba(new_patient_scaled)[0, 1]

label = "Diabetes" if pred_class == 1 else "No Diabetes"
print(f"Prediction: {label} (probability of diabetes = {pred_proba:.2f})")
```

Prediction: No Diabetes (probability of diabetes = 0.18)

Get the input from User

```
def get_user_input(feature_ranges):
    """Prompt the user for each feature, show its valid range, and validate the entry."""
    user_values = {}
    print("Enter the patient's details below.")
    print("(Valid ranges are shown based on the training data.)\n")
    for feature, (low, high) in feature_ranges.items():
        while True:
            raw = input(f"{feature} (range: {low:g} - {high:g}): ").strip()
            try:
                value = float(raw)
            except ValueError:
                print("Please enter a numeric value.")
                continue
```

```

        if value < low or value > high:
            print(f" Out of range. Please enter a value between {low:g} and {high:g}.")
            continue
        user_values[feature] = value
        break
    return user_values

# Build valid ranges directly from the cleaned training data (min, max per feature)
feature_ranges = {col: (df_clean[col].min(), df_clean[col].max()) for col in X.columns}

user_values = get_user_input(feature_ranges)
new_patient = pd.DataFrame([user_values])[X.columns] # keep the same column order as training

new_patient_scaled = scaler.transform(new_patient)
pred_class = knn_final.predict(new_patient_scaled)[0]
pred_proba = knn_final.predict_proba(new_patient_scaled)[0, 1]

label = "Diabetes" if pred_class == 1 else "No Diabetes"
print("\n--- Patient entered ---")
print(new_patient.to_string(index=False))
print(f"\nPrediction: {label} (probability of diabetes = {pred_proba:.2f})")

```

Enter the patient's details below.
(Valid ranges are shown based on the training data.)

```

Pregnancies (range: 0 - 17): 10
Glucose (range: 44 - 199): 78
BloodPressure (range: 24 - 122): 76
SkinThickness (range: 7 - 99): 89
Insulin (range: 14 - 846): 90
BMI (range: 18.2 - 67.1): 79.0
    Out of range. Please enter a value between 18.2 and 67.1.
BMI (range: 18.2 - 67.1): 19.0
DiabetesPedigreeFunction (range: 0.078 - 2.42): 2.1
Age (range: 21 - 81): 35

```

--- Patient entered ---

Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age
10.0	78.0	76.0	89.0	90.0	19.0	2.1	35.0

Prediction: Diabetes (probability of diabetes = 0.59)

Save Model

```

from sklearn.pipeline import Pipeline
import joblib

pipeline = Pipeline([
    ('scaler', scaler),
    ('knn', knn_final)
])

joblib.dump(pipeline, "KNN_Diabetes_Model.joblib")

['KNN_Diabetes_Model.joblib']

```

Flask Code :-

```

from flask import Flask, render_template, request
from joblib import load
import numpy as np

app = Flask(__name__)

# Load complete KNN pipeline
# Pipeline contains StandardScaler + KNN
model = load("KNN_Diabetes_Model.joblib")

@app.route('/')
def home():
    return render_template("index.html")

@app.route('/predict', methods=['POST'])

```

```
def predict():

    # -----
    # Get input values
    # -----
    pregnancies = float(request.form['pregnancies'])
    glucose = float(request.form['glucose'])
    blood_pressure = float(request.form['blood_pressure'])
    skin_thickness = float(request.form['skin_thickness'])
    insulin = float(request.form['insulin'])
    bmi = float(request.form['bmi'])
    diabetes_pedigree = float(request.form['diabetes_pedigree'])
    age = float(request.form['age'])

    # -----
    # Create input array
    # IMPORTANT:
    # Feature order must match
    # training data
    # -----
    input_data = np.array([[
        pregnancies,
        glucose,
        blood_pressure,
        skin_thickness,
        insulin,
        bmi,
        diabetes_pedigree,
        age
    ]])

    # -----
    # Prediction
    # Pipeline automatically:
    # 1. Scales the input
    # 2. Applies KNN
    # -----
    prediction = model.predict(input_data)

    # -----
    # Convert prediction to result
    # -----
    if prediction[0] == 1:
        result = "Diabetes"
    else:
        result = "No Diabetes"

    # -----
    # Send result to HTML
    # -----
    return render_template(
        "index.html",
        prediction_text=f"Prediction : {result}"
    )

if __name__ == "__main__":
    app.run(debug=True)

* Serving Flask app '__main__'
* Debug mode: on
INFO:werkzeug:WARNING: This is a development server. Do not use it in a production deployment. Use a production WSGI server inst
* Running on http://127.0.0.1:5000
INFO:werkzeug:Press CTRL+C to quit
INFO:werkzeug: * Restarting with watchdog (inotify)
```

HTML Code :-

```
<!DOCTYPE html>
<html lang="en">

<head>
```

```
<meta charset="UTF-8">

<meta name="viewport" content="width=device-width, initial-scale=1.0">

<title>KNN Diabetes Classification</title>

<style>

  *{
    box-sizing: border-box;
  }

  body {
    margin: 0;
    font-family: Arial, sans-serif;
    background: #f4f7fb;
    color: #333;
  }

  .container {
    width: 95%;
    max-width: 1000px;
    margin: 30px auto;
  }

  .card {
    background: white;
    padding: 30px;
    border-radius: 12px;
    box-shadow: 0 4px 15px rgba(0, 0, 0, 0.08);
  }

  h1 {
    text-align: center;
    color: #1769aa;
    margin-bottom: 8px;
  }

  .subtitle {
    text-align: center;
    color: #666;
    margin-bottom: 10px;
  }

  .instruction {
    text-align: center;
    color: #777;
    font-size: 13px;
    margin-bottom: 25px;
  }

  .section-title {
    background: #eaf3fb;
    color: #1769aa;
    padding: 12px 15px;
    border-radius: 6px;
    margin-top: 25px;
    margin-bottom: 15px;
    font-size: 18px;
    font-weight: bold;
  }

  .section-note {
    font-size: 13px;
    color: #777;
    margin-bottom: 12px;
  }

  .input-grid {
    display: grid;
    grid-template-columns: repeat(2, 1fr);
    gap: 15px;
  }

  .input-group {
    display: flex;
```

```

    flex-direction: column;
  }

  label {
    font-size: 14px;
    margin-bottom: 6px;
    font-weight: bold;
  }

  input {
    padding: 10px;
    border: 1px solid #ccc;
    border-radius: 5px;
    font-size: 14px;
  }

  input:focus {
    outline: none;
    border-color: #1769aa;
  }

  input::placeholder {
    color: #999;
    font-size: 12px;
  }

  .button-area {
    text-align: center;
    margin-top: 30px;
  }

  button {
    background: #1769aa;
    color: white;
    border: none;
    padding: 13px 35px;
    border-radius: 6px;
    font-size: 16px;
    cursor: pointer;
  }

  button:hover {
    background: #0d568d;
  }

  .result {
    margin-top: 25px;
    padding: 15px;
    text-align: center;
    background: #e8f7ed;
    color: #218838;
    border-radius: 6px;
    font-size: 18px;
    font-weight: bold;
  }

  .probability {
    margin-top: 10px;
    padding: 12px;
    text-align: center;
    background: #eef5ff;
    color: #1769aa;
    border-radius: 6px;
    font-size: 15px;
    font-weight: bold;
  }

  .footer {
    text-align: center;
    color: #777;
    margin-top: 20px;
    font-size: 13px;
  }

  @media(max-width: 700px) {

    .input-grid {
      grid-template-columns: 1fr;

```

```

    }

    .card {
      padding: 20px;
    }

  }

</style>
</head>

<body>

<div class="container">

  <div class="card">

    <h1>KNN Diabetes Classification</h1>

    <div class="subtitle">
      Classification using K-Nearest Neighbors
    </div>

    <div class="instruction">
      Enter the patient's 8 medical details below.
    </div>

    <form action="/predict" method="POST">

      <!-- ===== -->
      <!-- PATIENT FEATURES -->
      <!-- ===== -->

      <div class="section-title">
        Patient Medical Features
      </div>

      <div class="section-note">
        Enter numerical values for all the features used by the KNN classifier.
      </div>

      <div class="input-grid">

        <!-- Pregnancies -->

        <div class="input-group">

          <label>Pregnancies</label>

          <input type="number"
            step="any"
            name="pregnancies"
            placeholder="Example: 2"
            min="0"
            required>

        </div>

        <!-- Glucose -->

        <div class="input-group">

          <label>Glucose</label>

          <input type="number"
            step="any"
            name="glucose"
            placeholder="Example: 130"
            min="0"
            required>

```

```

</div>

<!-- Blood Pressure -->

<div class="input-group">

    <label>Blood Pressure</label>

    <input type="number"
          step="any"
          name="blood_pressure"
          placeholder="Example: 70"
          min="0"
          required>

</div>

<!-- Skin Thickness -->

<div class="input-group">

    <label>Skin Thickness</label>

    <input type="number"
          step="any"
          name="skin_thickness"
          placeholder="Example: 25"
          min="0"
          required>

</div>

<!-- Insulin -->

<div class="input-group">

    <label>Insulin</label>

    <input type="number"
          step="any"
          name="insulin"
          placeholder="Example: 100"
          min="0"
          required>

</div>

<!-- BMI -->

<div class="input-group">

    <label>BMI</label>

    <input type="number"
          step="any"
          name="bmi"
          placeholder="Example: 31.5"
          min="0"
          required>

</div>

<!-- Diabetes Pedigree Function -->

<div class="input-group">

    <label>Diabetes Pedigree Function</label>

    <input type="number"
          step="any"
          name="diabetes_pedigree"
          placeholder="Example: 0.45"
          min="0"
          required>

</div>

```

```
min= 0
required>

</div>

<!-- Age -->

<div class="input-group">

  <label>Age</label>

  <input type="number"
        step="any"
        name="age"
        placeholder="Example: 35"
        min="0"
        required>

</div>

</div>

<!-- ===== -->
<!-- PREDICT BUTTON -->
<!-- ===== -->

<div class="button-area">

  <button type="submit">
```